

Beyond Pairwise GUI Quality Judgments: Multi-Dimensional Evaluation with Modern Multimodal LLMs

Xinyang Xie

Department of Informatics, Technische Universitaet Clausthal,
Clausthal-Zellerfeld, Germany.

Abstract

Assessing the quality of graphical user interfaces (GUIs) is becoming increasingly important as large language models are used to generate interface designs and frontend code. UIClip showed that vision-language models can make pairwise quality judgments from screenshots and textual descriptions, but the original setting is limited by older models, binary decisions, and simple prompting. This proposal revisits that problem with current multimodal large language models (MLLMs). The study combines pairwise comparison with rubric-based assessment across layout quality, visual hierarchy, information organization, and perceived usability. Methodologically, it focuses on a screenshot-based evaluation set, a human annotation protocol with agreement analysis, and a comparison of judging strategies such as zero-shot prompting, rubric-guided scoring, repeated judging, and **order-swapped pairwise prompts**. The goal is to understand which evaluation setup best aligns model judgments with human judgments and which failure patterns remain in GUI quality assessment.



1 Introduction and Motivation

The quality of a graphical user interface (GUI) shapes **how efficiently users complete tasks** and how they perceive a software system. As generative models are increasingly used to produce interface mockups and frontend code, the question of how UI quality can be evaluated becomes more pressing.

UIClip [1] provided an important first step by showing that vision-language models can perform pairwise comparisons of UI screenshots and identify which of two interfaces appears better. At the same time, the original setup leaves several open

questions: it relies on older models, reduces quality assessment to binary comparisons, and does not examine how newer LLM-as-a-judge strategies affect reliability.

Recent work on improving UI generation from designer feedback [2] suggests that high-quality, workflow-aligned feedback can be more informative than large amounts of coarse ranking data. In parallel, the broader LLM-as-a-judge literature [3] has introduced prompting, consistency, and aggregation strategies that may also be useful for GUI evaluation. Taken together, these developments motivate a more methodologically explicit re-evaluation of GUI quality judgment with contemporary multimodal models.

This thesis therefore asks whether modern MLLMs can approximate human judgments of GUI quality beyond pairwise comparison, and how such judgments should be elicited and evaluated in a reliable way.

2 Planned Approach

The planned work centers on an evaluation pipeline with four main components: dataset construction, human annotation, rubric design, and the comparison of different LLM-as-a-judge strategies. 

2.1 Datasets and Evaluation Material

The study will build primarily on publicly available [screenshot-based resources](#) from UIClip. To avoid evaluating only older interfaces, this material will be complemented, if needed, by a smaller curated set of contemporary web or mobile screenshots. The resulting evaluation material will consist of two subsets: a pairwise comparison subset for overall preference judgments and a single-image subset for structured multi-dimensional assessment.

2.2 Human Annotation Protocol

Human judgments will be collected for a selected evaluation subset. For pairwise tasks, annotators will compare two GUI screenshots and indicate which interface is better overall, with an optional tie label if required. For single-image tasks, annotators will score each GUI with a small rubric and provide brief justifications. To improve the reliability of the reference labels, each item will be annotated by multiple raters, and inter-rater reliability will be measured, for example with Krippendorff's alpha.

2.3 Multi-dimensional GUI Rubric

To move beyond binary comparison, I will define a small rubric grounded in common visual design principles and usability heuristics. The initial dimensions are:

- **Layout quality:** spatial organization, alignment, spacing, and overall balance.
- **Visual hierarchy:** whether contrast, size, and color guide attention appropriately.
- **Information organization:** clarity of grouping and logical ordering of content.
- **Perceived usability:** how easily key actions and interface purpose can be inferred from a static screenshot.



Each dimension will be scored on a fixed ordinal scale with short written anchors so that the same rubric can be applied by both human annotators and model judges. The perceived usability dimension is intentionally limited to what can reasonably be inferred from a static screen, rather than to full interaction quality.

2.4 Prompting and Judging Strategies

Several judging strategies will be compared. A zero-shot baseline will ask the model directly for a pairwise preference or a rubric-based score. A few-shot variant will add one or two worked examples to test whether light calibration improves judgment quality. A rubric-guided condition will provide explicit scoring instructions and require a structured output format. To test robustness, the same prompt will be repeated to measure self-consistency. For pairwise tasks, image order will be swapped to test positional bias. In addition, the study will compare single-model judgments with simple aggregation strategies such as majority voting for pairwise tasks and averaged scores for rubric-based tasks.



3 Evaluation

Several current multimodal models will be tested as judges, including representative models from the GPT, Gemini, and Claude families and, where feasible, at least one strong open multimodal model. The pairwise setting will be evaluated through agreement with human majority labels, while the rubric-based setting will be evaluated through rank correlation and score differences between model outputs and human ratings.



Beyond task accuracy, the evaluation will explicitly examine reliability. This includes inter-rater agreement in the human annotations, consistency across repeated model judgments, stability under pairwise order reversal, and agreement across different models. Computational cost will also be tracked in order to assess practical trade-offs between evaluation quality and API usage.

Finally, a qualitative error analysis will examine typical failure cases, such as cluttered layouts, minimal interfaces, text-heavy screens, and visually appealing but less usable designs. Together, these analyses should clarify how well modern MLLMs support UIClip-style evaluation, which rubric dimensions are most reliable, and which judging strategies align best with human assessment.

References

- [1] Wu, J., Peng, Y.-H., Li, X.Y.A., *et al.*: UIClip: A data-driven model for assessing user interface design. In: Proc. ACM UIST (2024)
- [2] Wu, J., Swearngin, A., Vajjala, A.K., *et al.*: Improving user interface generation models from designer feedback. In: Proc. CHI (2026)
- [3] Zheng, L., Chiang, W.-L., Sheng, Y., *et al.*: Judging LLM-as-a-judge with MT-bench and chatbot arena. In: NeurIPS (2023)